

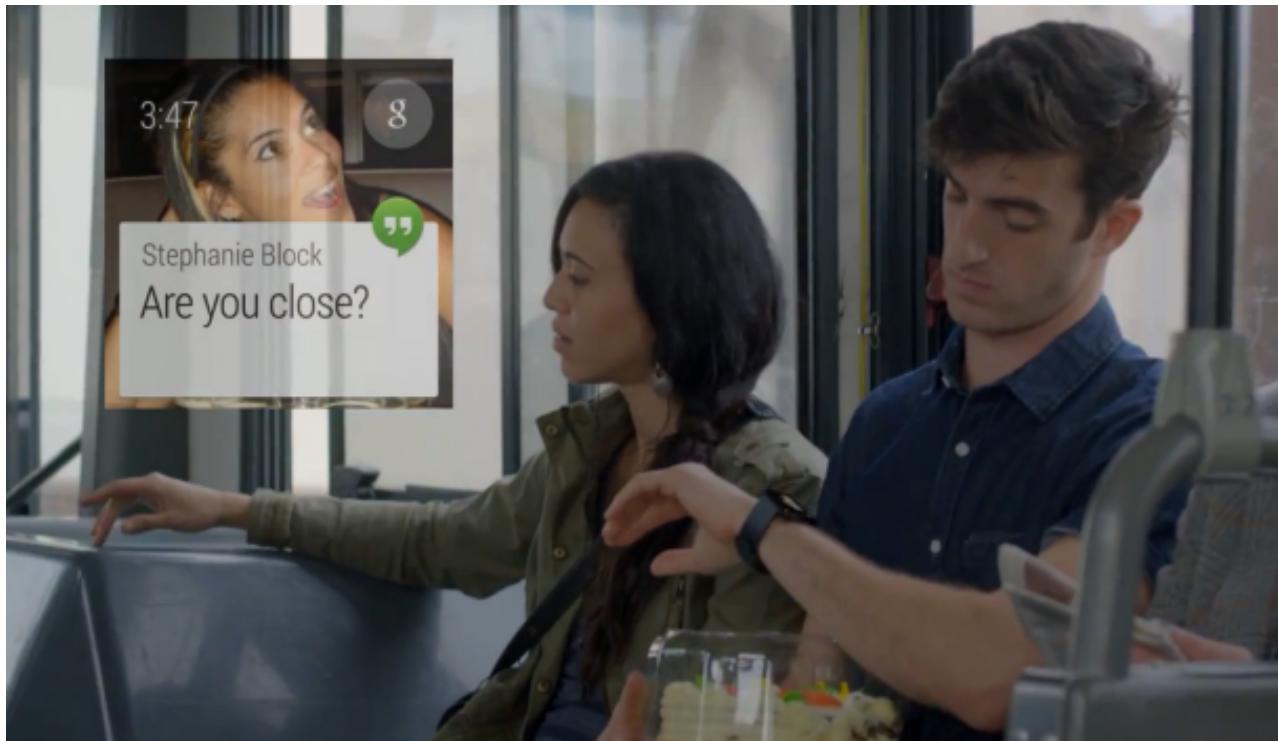
Understanding the Characteristics of Android Wear OS

Renju Liu and Felix Xiaozhu Lin

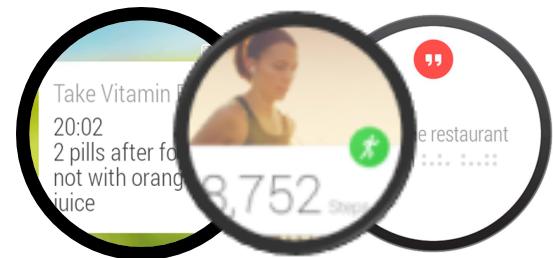
Purdue ECE

PURDUE
ENGINEERING





The Wearable stack



Top questions

- Wearables should enjoy
 - Baremetal performance
 - Baremetal efficiency
- In this talk: **Android Wear**
 - Are we close to baremetal?
 - What is going on *inside*?
 - How should the OS evolve?

Observation -- Symptoms

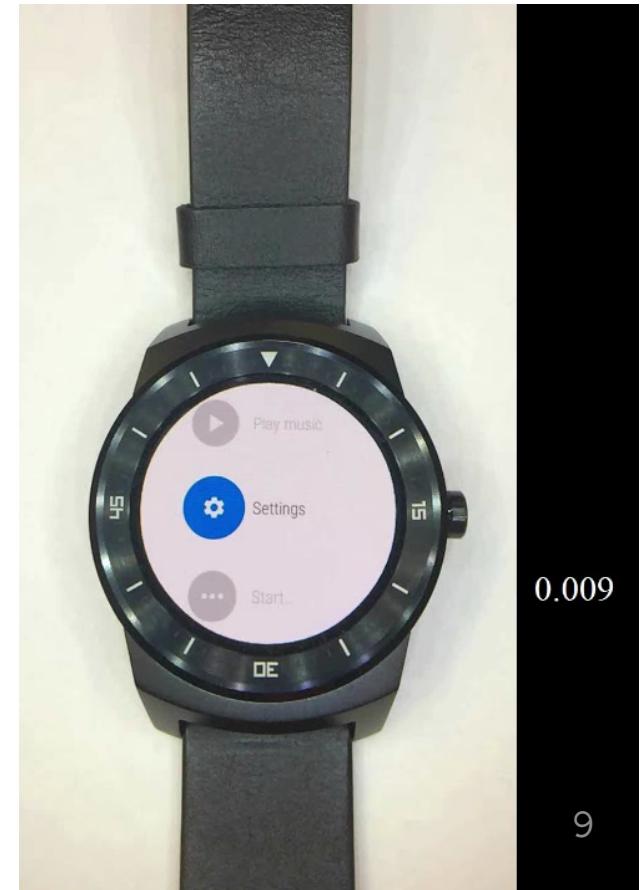
- The current performance & efficiency are far from baremetal
- **Pacing – inefficient**
 - face update: 400ms 88% busy

Clock face update

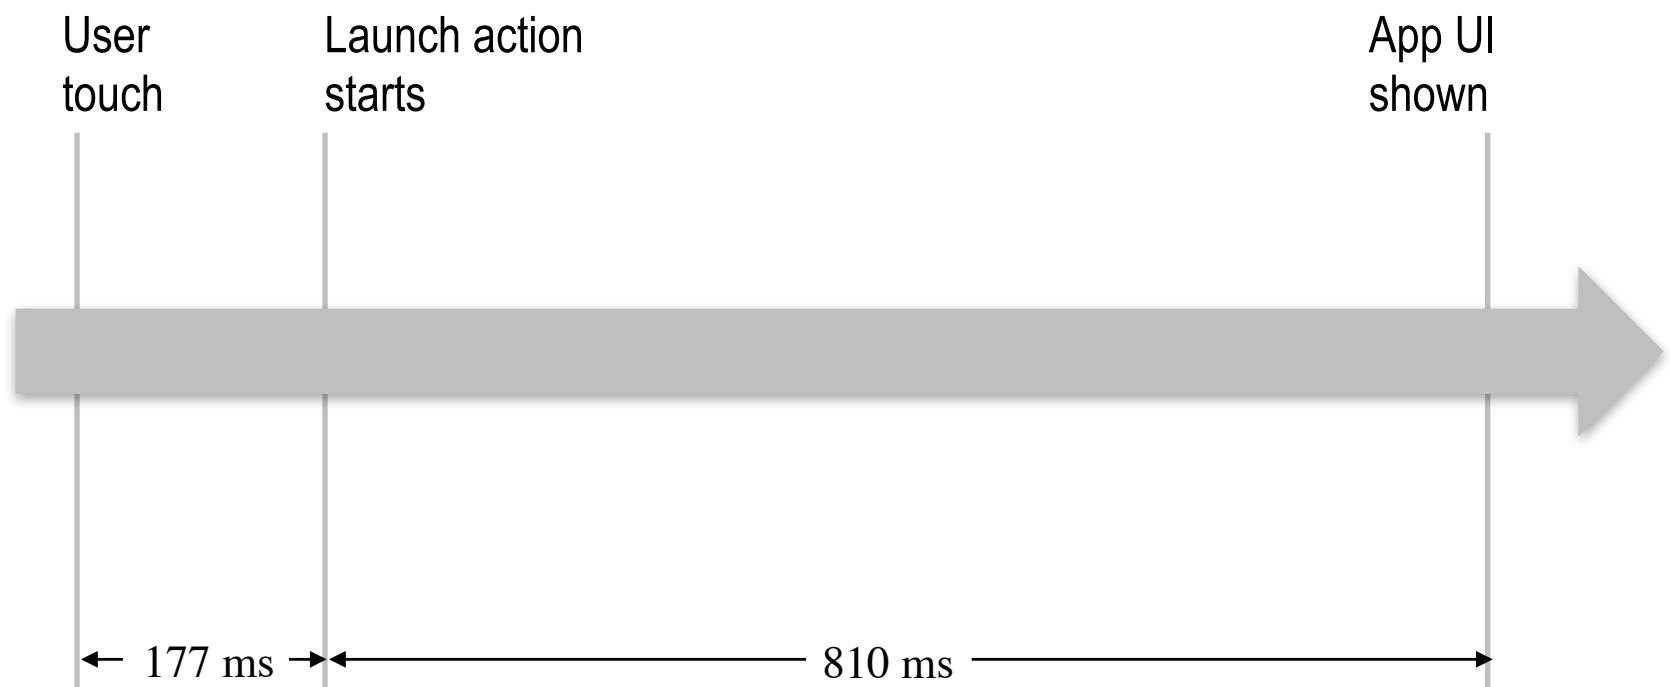


Observation -- Symptoms

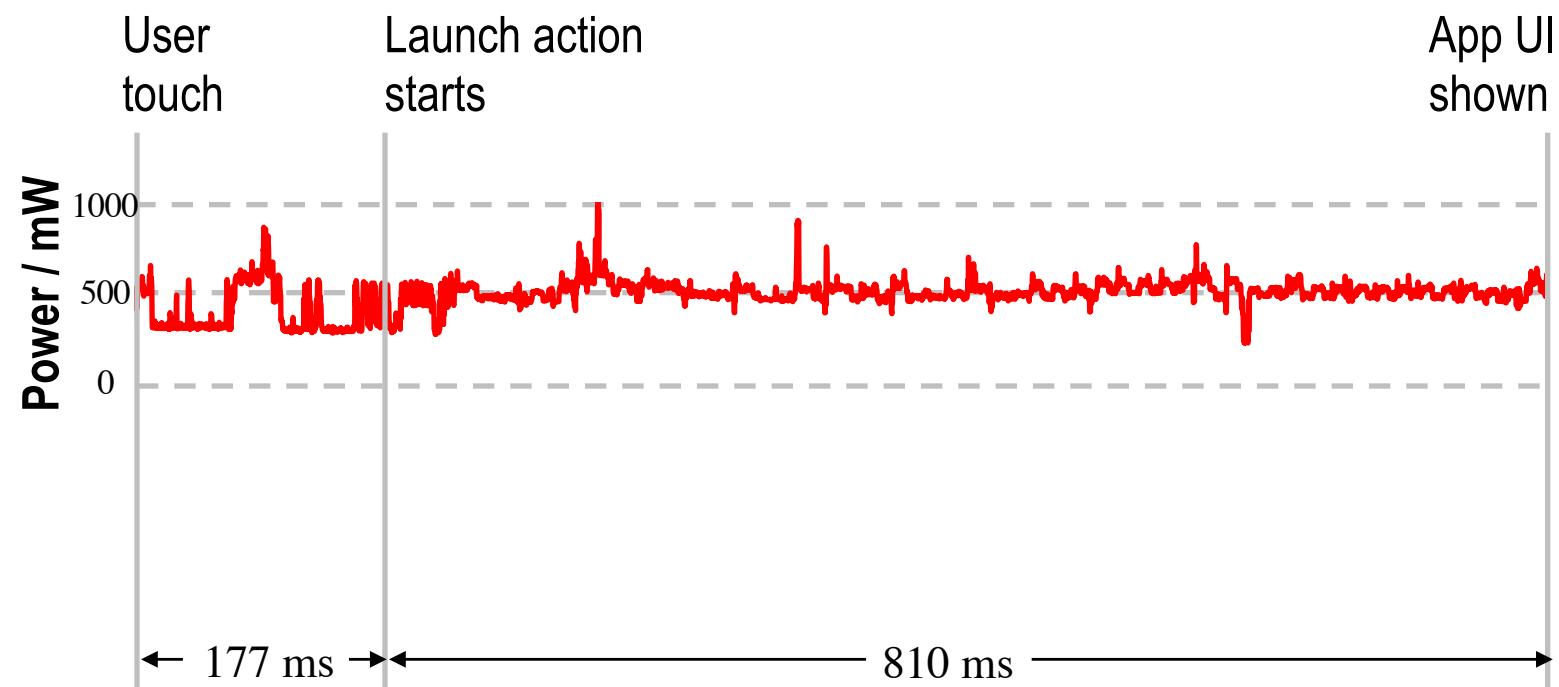
- The current performance & efficiency are far from baremetal Launch “settings”
- **Pacing – inefficient**
 - face update: 400ms 88% busy
- **Racing – slow**
 - Launch an in-mem app: 1 sec



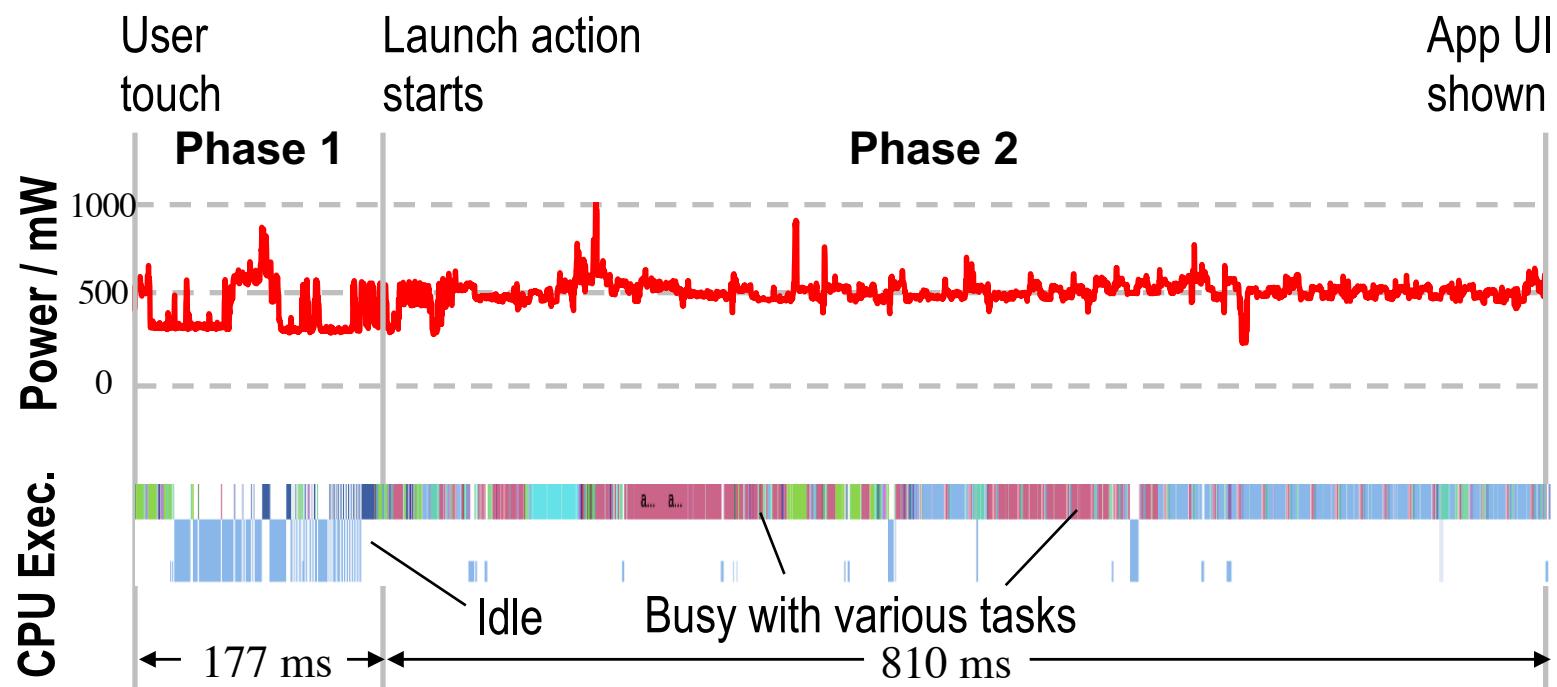
What happens underneath?



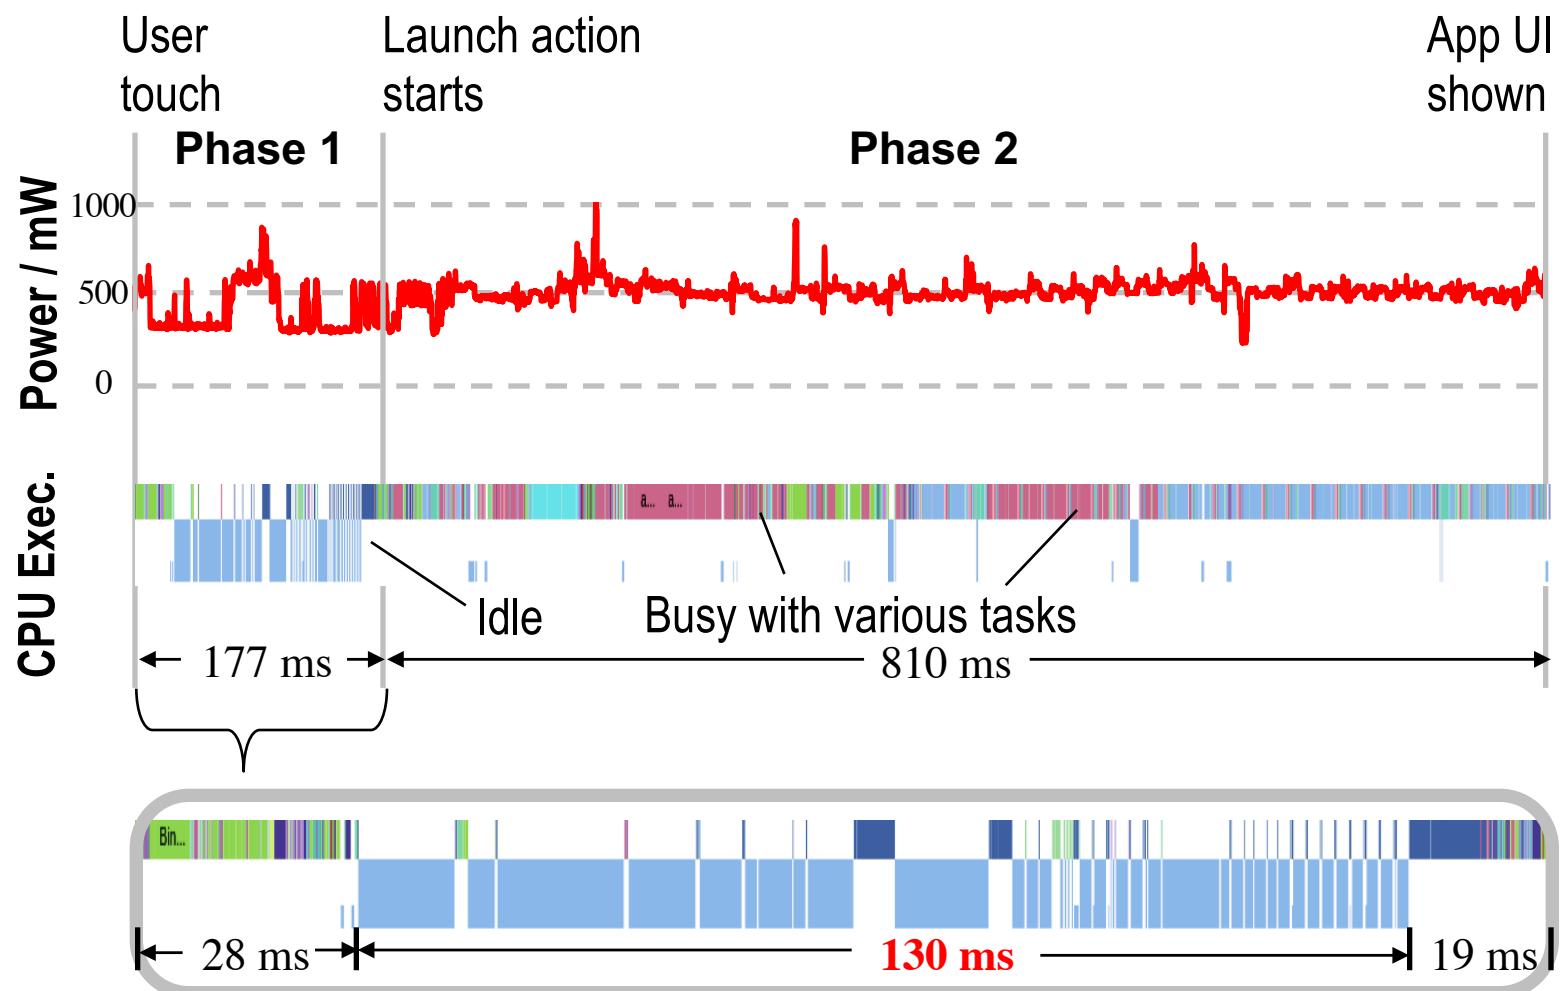
What happens underneath?



What happens underneath?



What happens underneath?



Four Aspects

CPU busy?

CPU idle?

Thread-level parallelism (TLP)

Microarchitectural behaviors

Won't talk about our methodologies

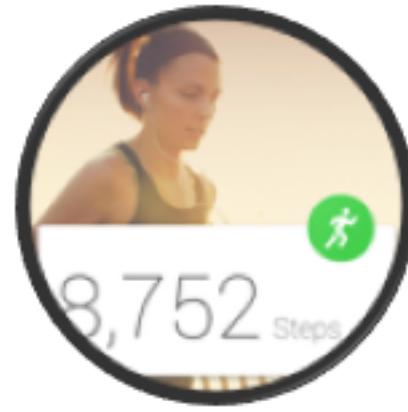
Profiling – Core Use Scenarios



Wakeup
Update
notification
wrist...



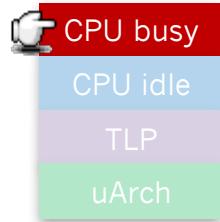
Single Input
launch apps
palming
voice...



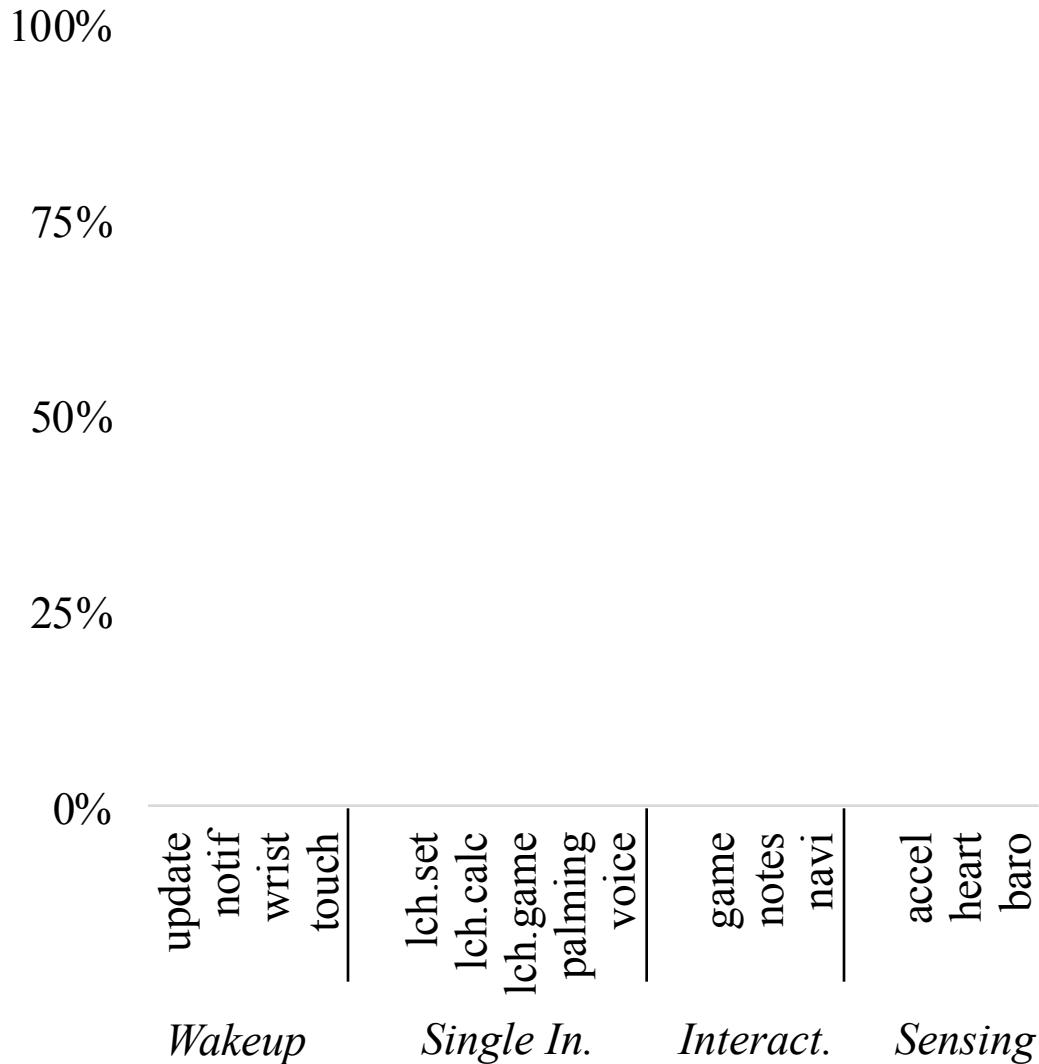
Sensing
Accel
heart
baro

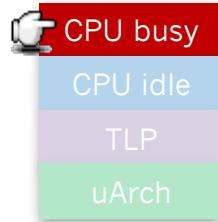


Interaction
Game
notes
navigation

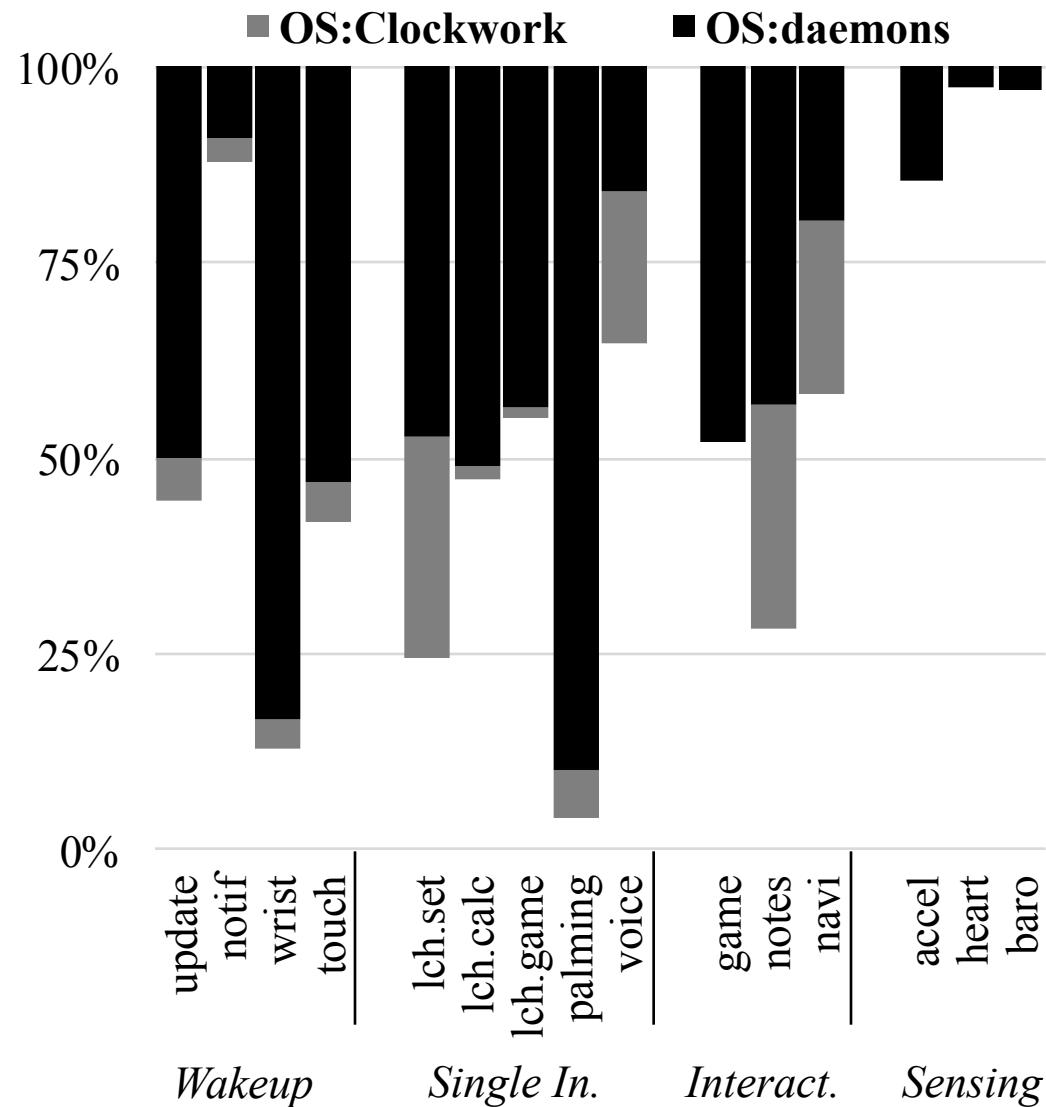


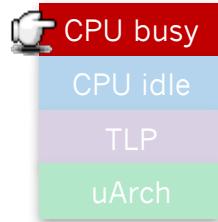
OS execution dominates CPU usage.



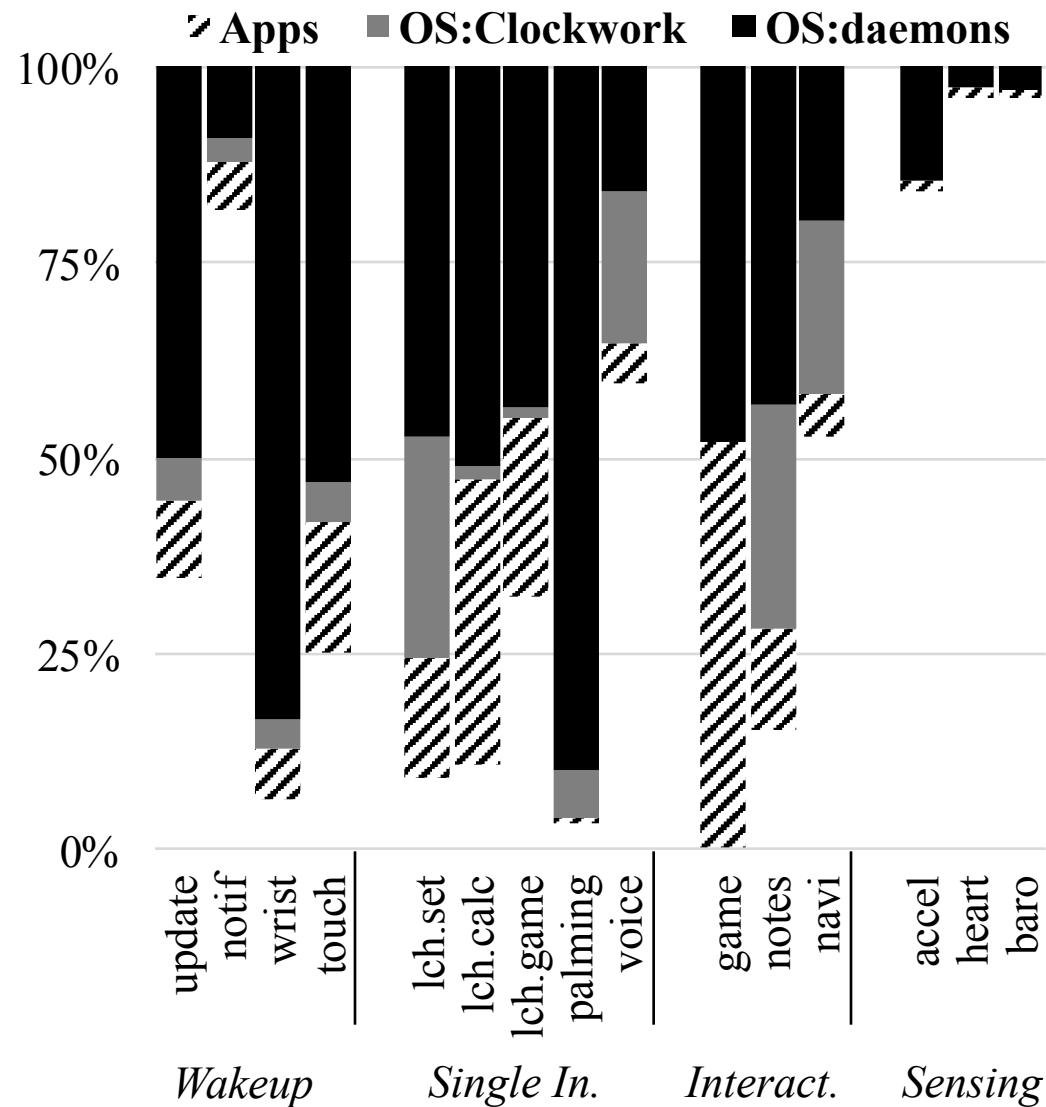


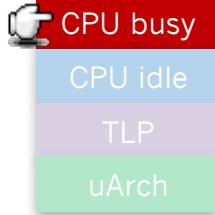
OS execution dominates CPU usage.



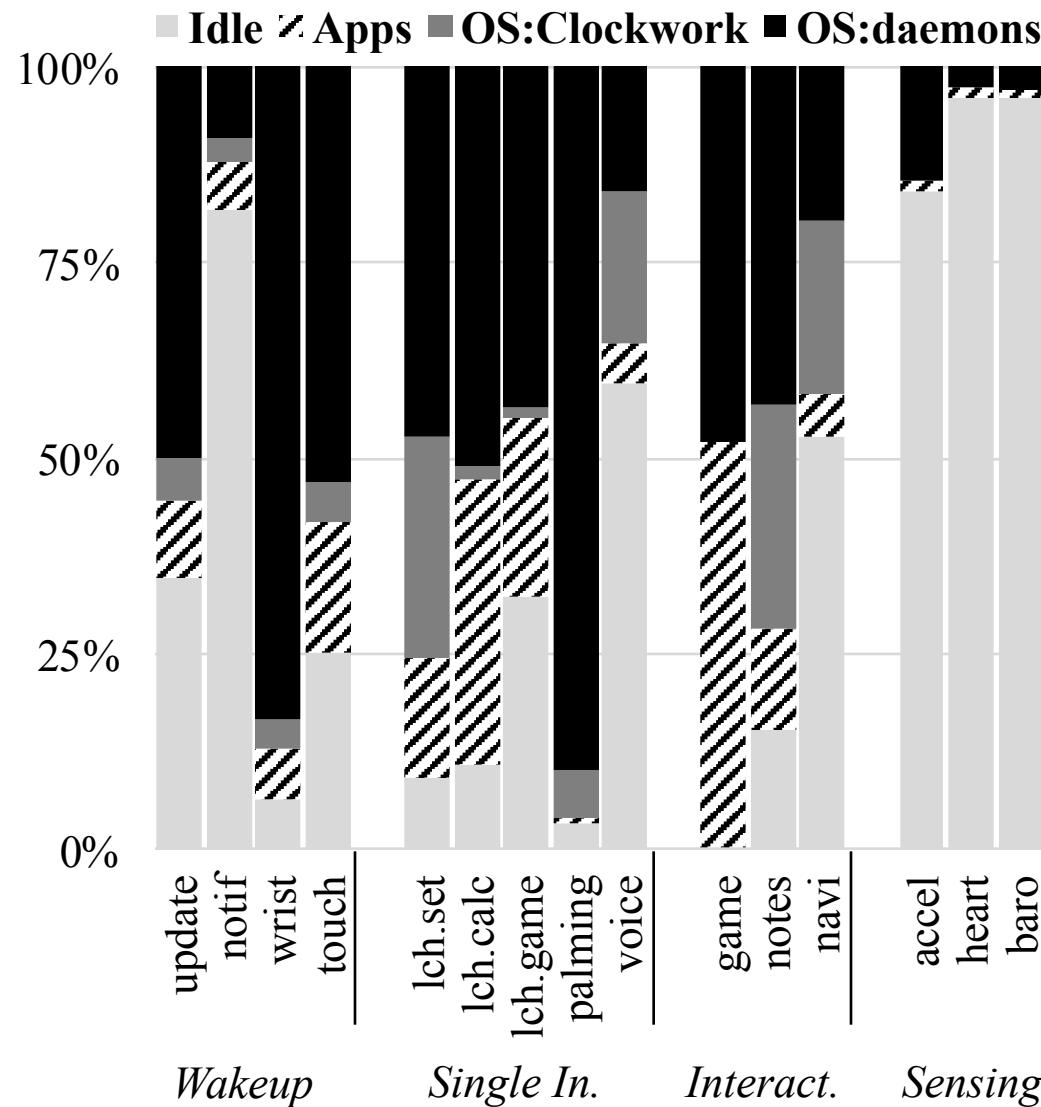


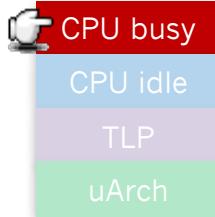
OS execution dominates CPU usage.



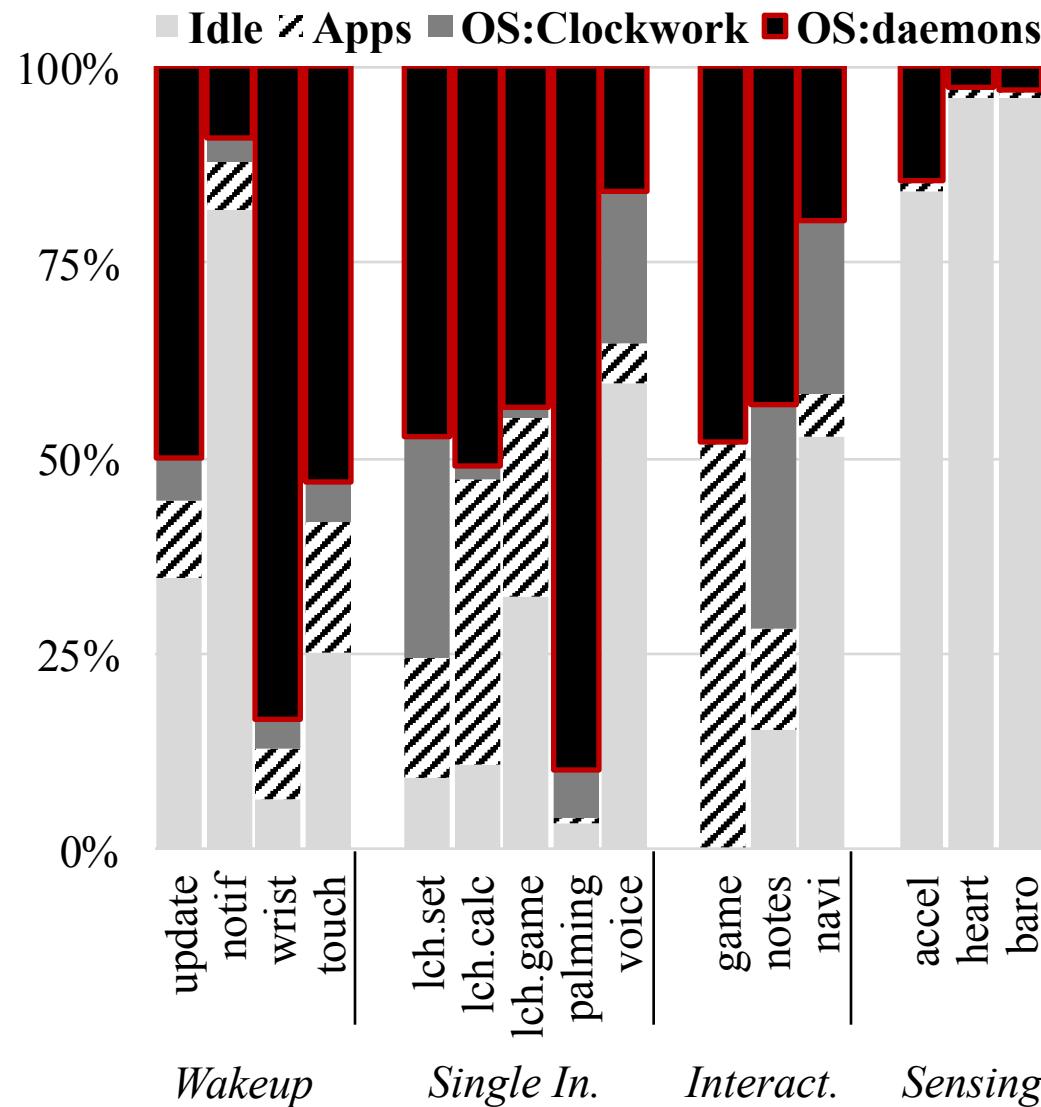


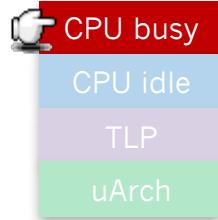
OS execution dominates CPU usage.



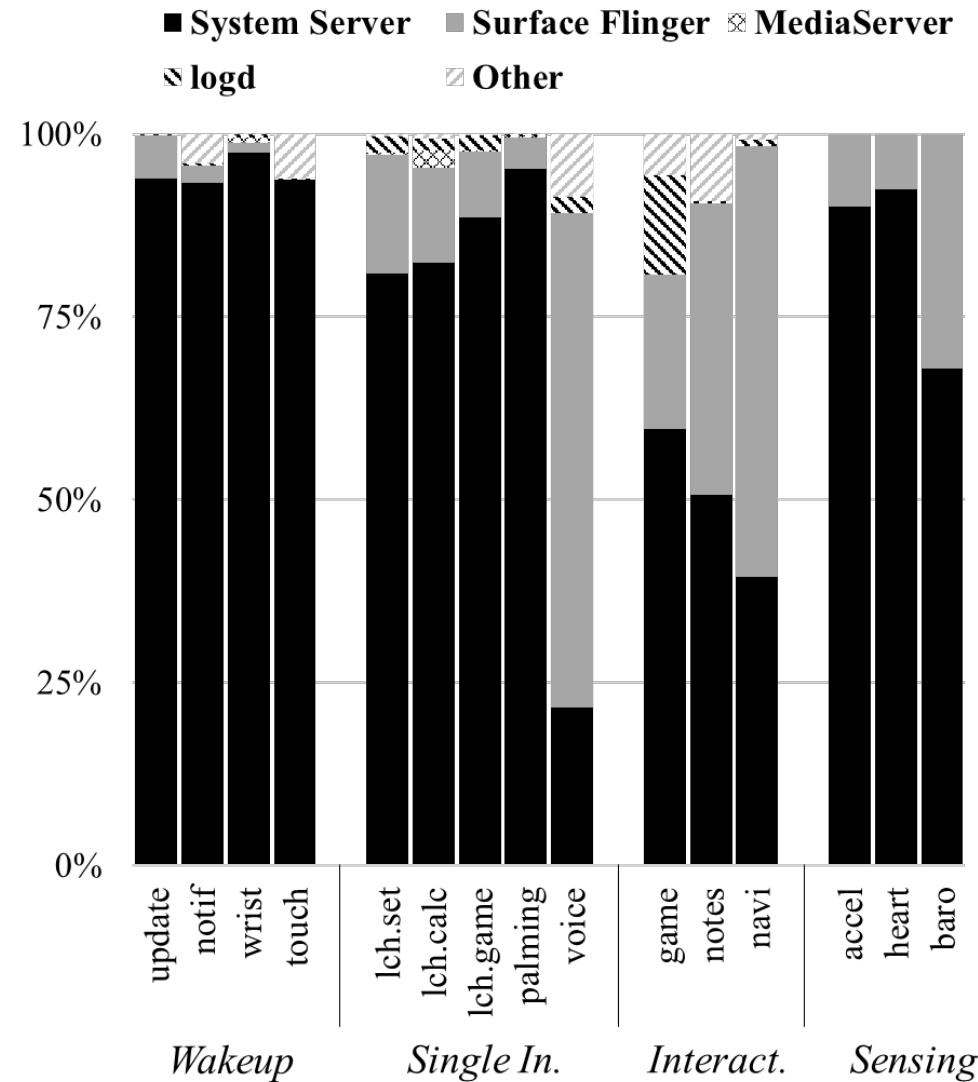


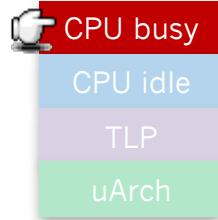
OS execution dominates CPU usage.



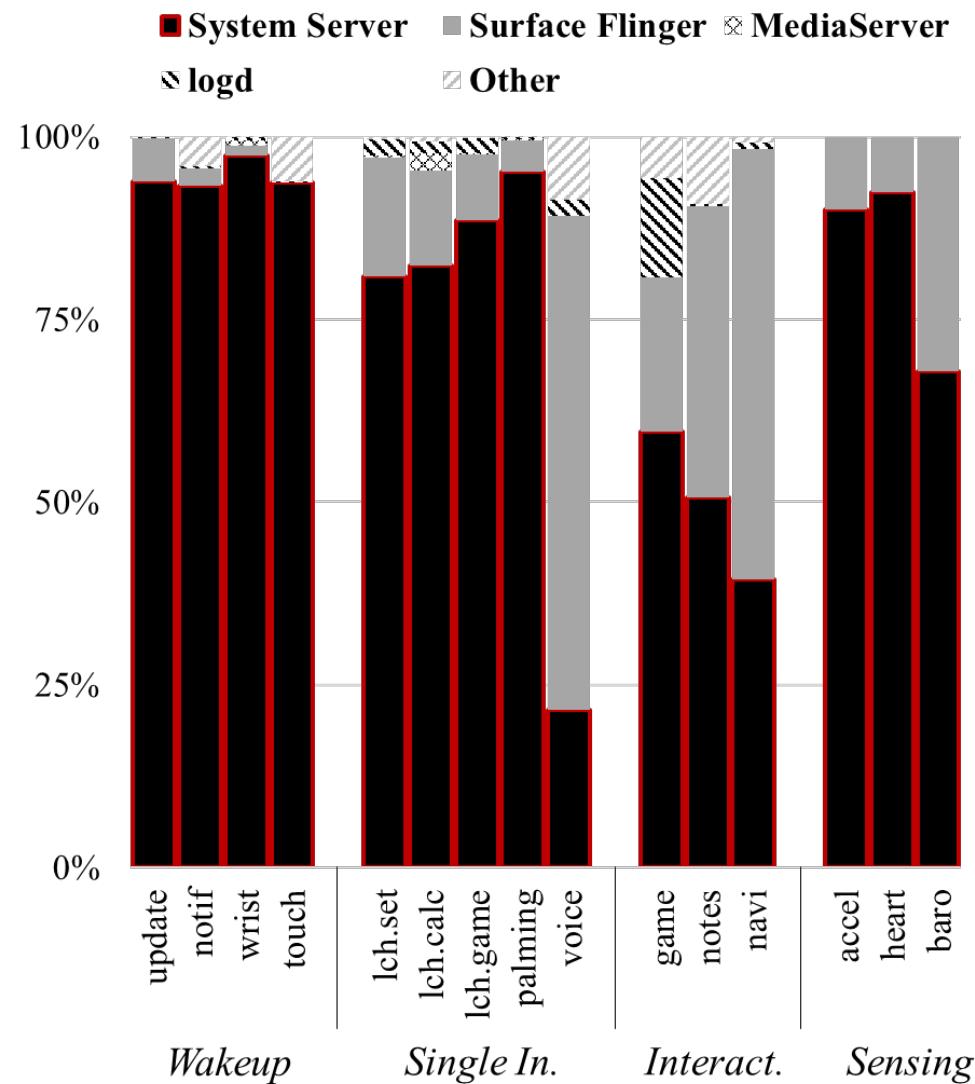


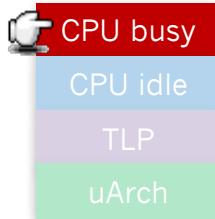
OS execution dominates CPU usage.



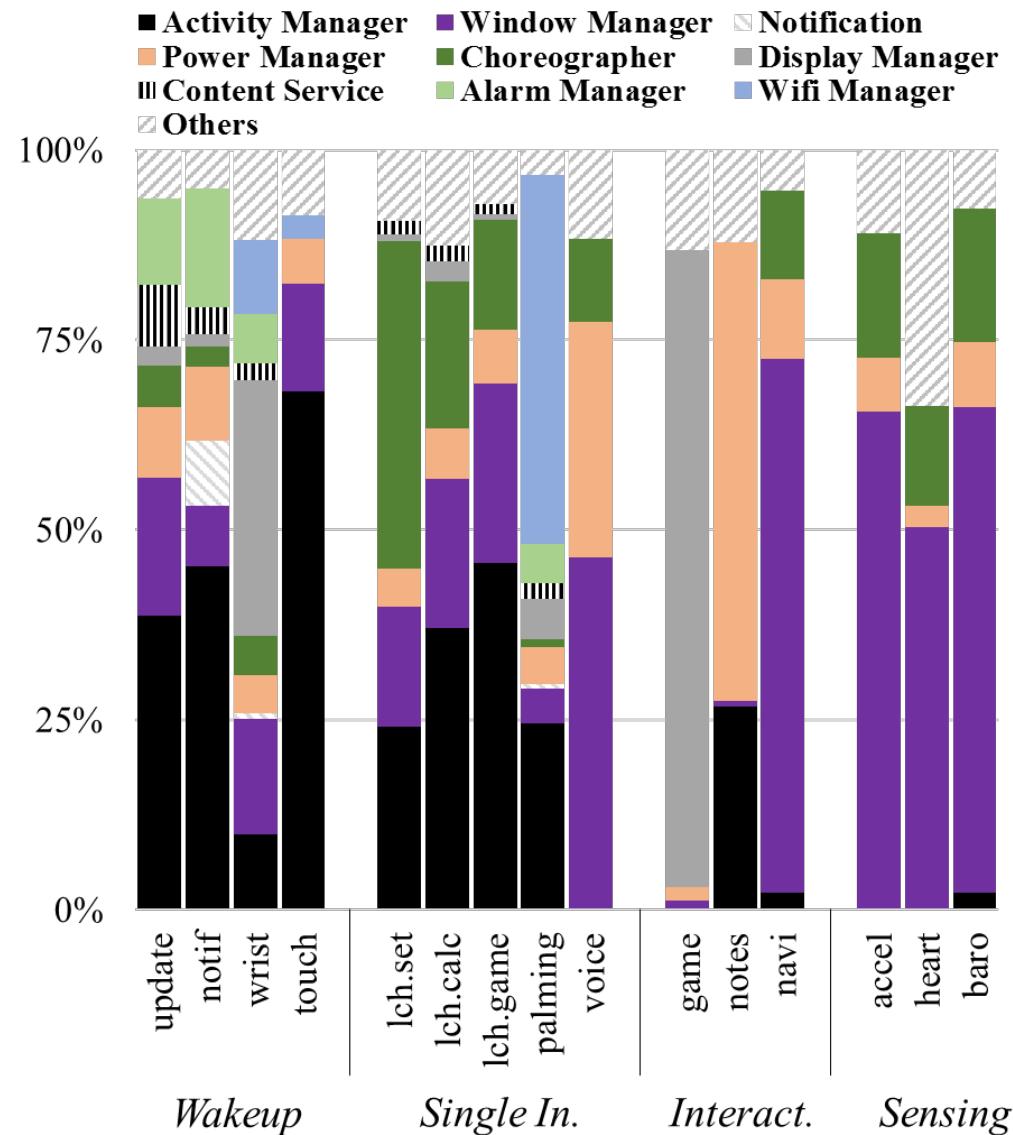


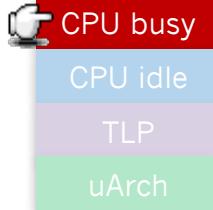
OS execution dominates CPU usage.



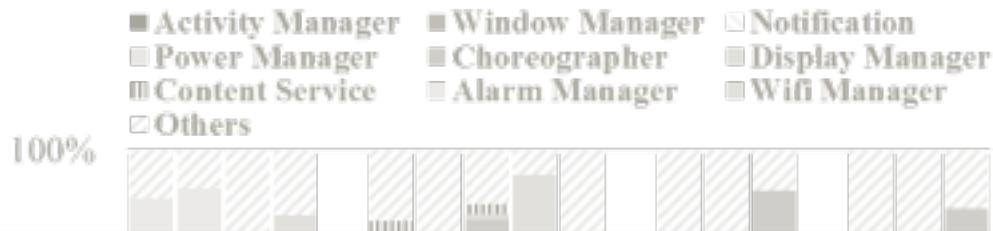


Costly OS services are ...





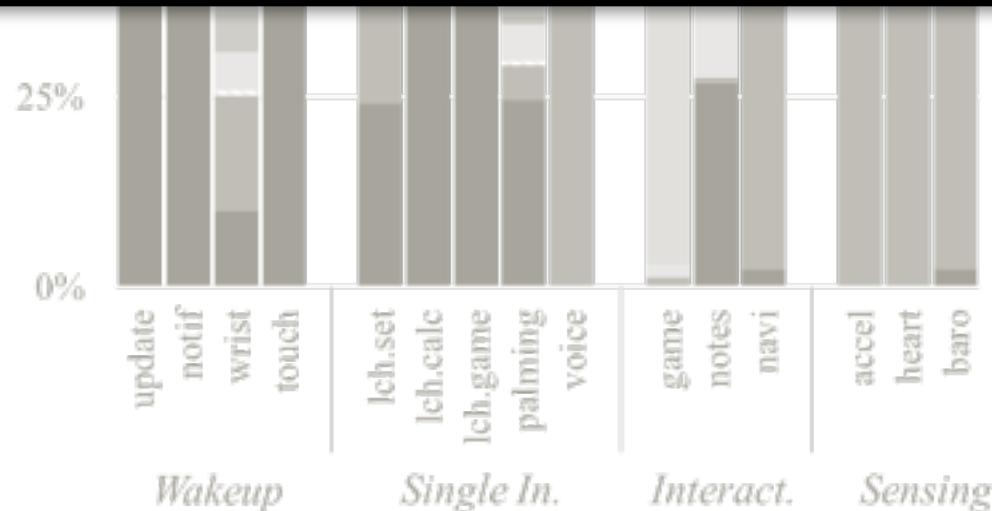
Costly OS services are likely cruft.

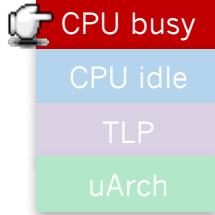


cruft  (krūft)

n.

1. Trash, debris, or other unwanted matter that accumulates over time.
2. Unnecessary digital information that accumulates over time, such as unneeded files or obsolete lines of code in software: *"By removing cruft, you can recover valuable disk space ... and reduce the chance of software conflicts"* (Joe Kissell).

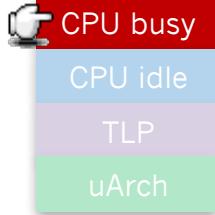




Hot functions: highly skewed distribution

Top 5 → >20% CPU cycles

Top 50 → >50% CPU cycles



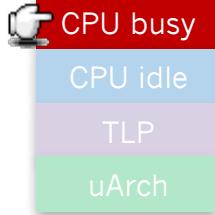
Hot functions: highly skewed distribution

Top 5 → >20% CPU cycles

Top 50 → >50% CPU cycles

Manipulating basic data structures

Legacy/improper OS designs



Hot functions: highly skewed distribution

Top 5 → >20% CPU cycles

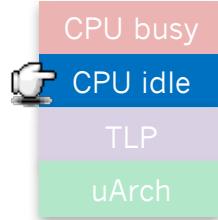
Top 50 → >50% CPU cycles

Manipulating basic data structures

Legacy/improper OS designs

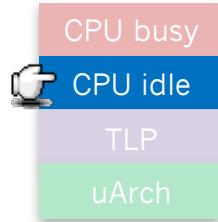
Anecdotes

Backlight UI layout low-mem killer



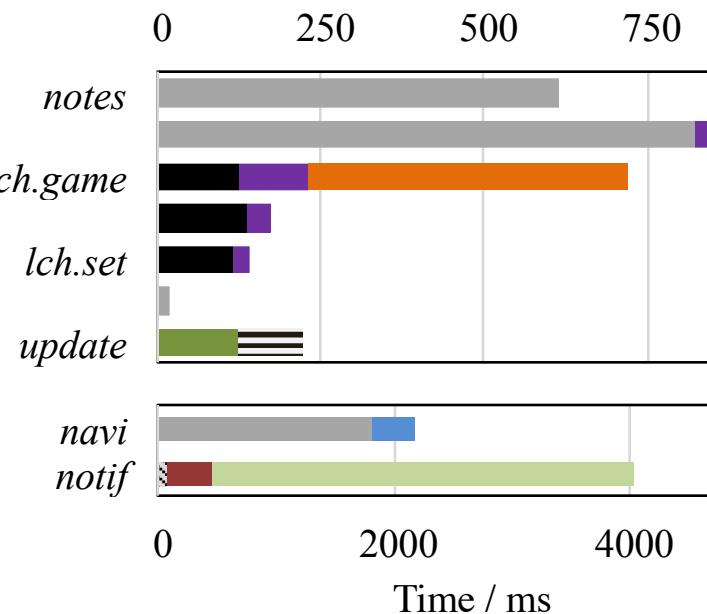
Idle episodes: plentiful and of various lengths

Time (ms)	Pct. Overall	Episodes	Pct. Explained	
614.1	17.1%	376	100.0%	<i>notes</i>
843.3	50.5%	352	100.0%	<i>voice</i>
722.6	50.9%	205	99.9%	<i>lch.game</i>
185.2	25.6%	110	92.9%	<i>lch.calc</i>
153.6	15.6%	120	91.4%	<i>lch.set</i>
16.8	10.6%	6	100.0%	<i>touch</i>
223.0	61.2%	44	100.0%	<i>update</i>
2173.0	52.80%	912	100.0%	<i>navi</i>
4035.6	86.80%	277	100.0%	<i>notif</i>

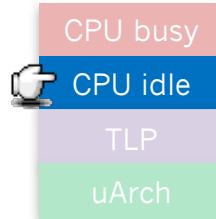


Idle anomalies are caused by ...

- Device suspend
- Voice UI
- Cont. interaction
- Cont. interact.+NetI/O
- Storage I/O
- User think
- Bluetooth tail time
- = OS shell policy
- App policy

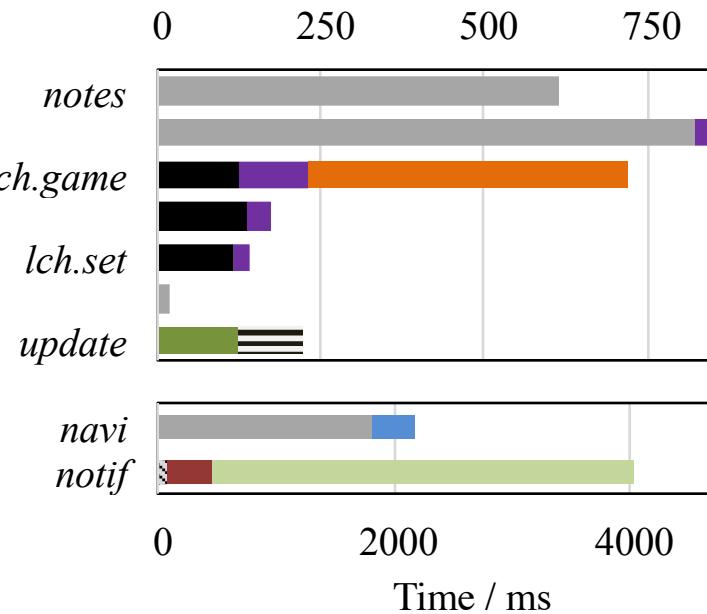


	Time (ms)	Pct. Overall	Episodes	Pct. Explained
notes	614.1	17.1%	376	100.0%
voice	843.3	50.5%	352	100.0%
lch.game	722.6	50.9%	205	99.9%
lch.calc	185.2	25.6%	110	92.9%
lch.set	153.6	15.6%	120	91.4%
touch	16.8	10.6%	6	100.0%
update	223.0	61.2%	44	100.0%
navi	2173.0	52.80%	912	100.0%
notif	4035.6	86.80%	277	100.0%



Idle anomalies are caused by ...

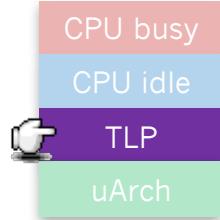
- Device suspend
- Voice UI
- Cont. interaction
- Cont. interact.+NetI/O
- Storage I/O
- User think
- Bluetooth tail time
- OS shell policy
- App policy



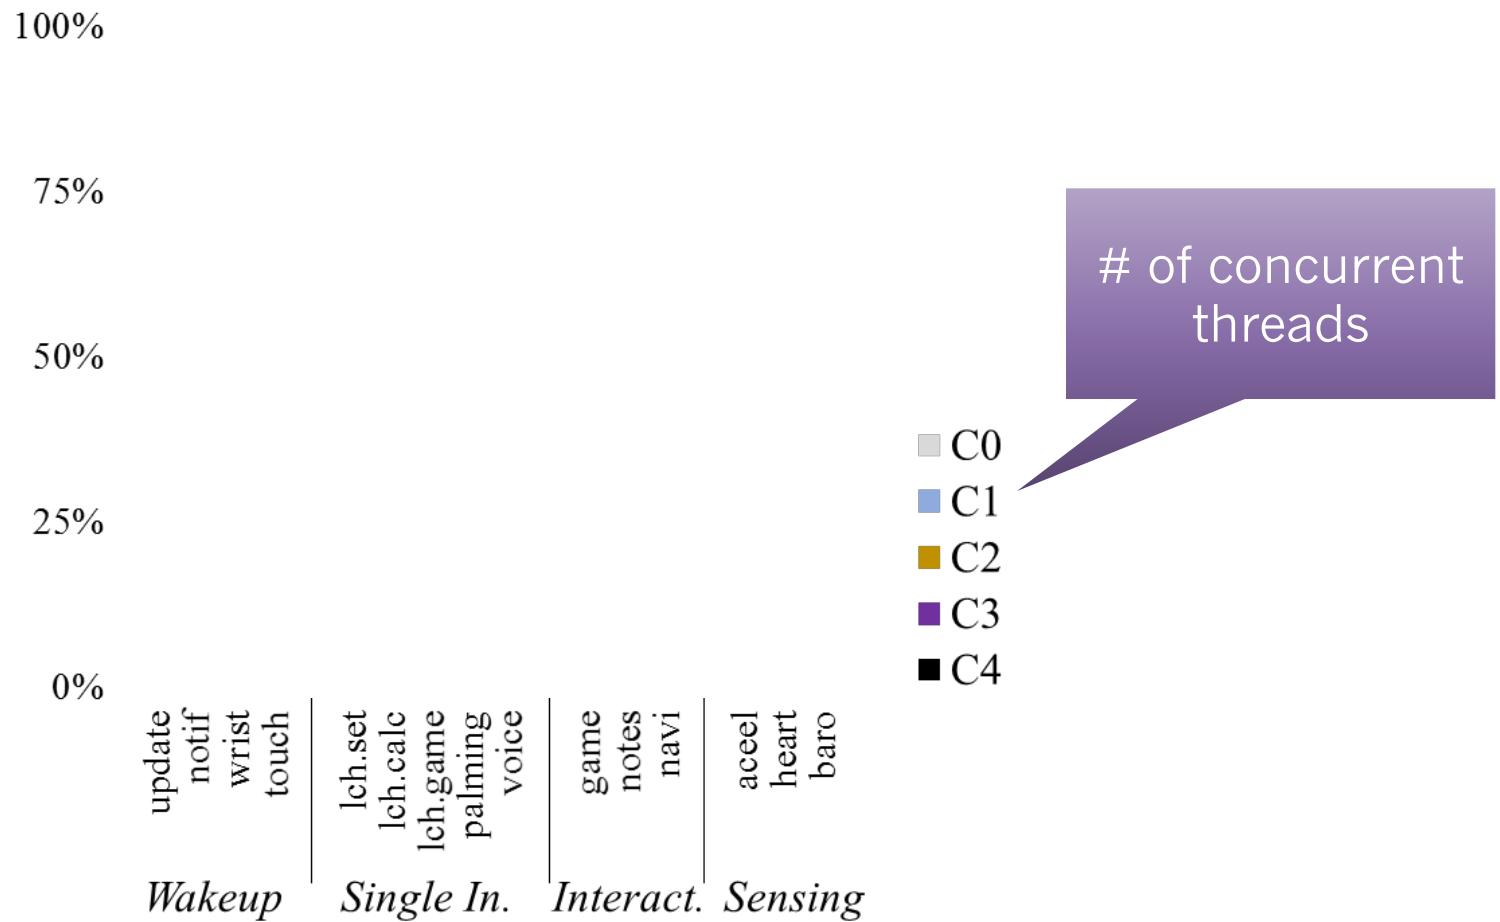
Time (ms)	Pct. Overall	Episodes	Pct. Explained
614.1	17.1%	376	100.0%
843.3	50.5%	352	100.0%
722.6	50.9%	205	99.9%
185.2	25.6%	110	92.9%
153.6	15.6%	120	91.4%
16.8	10.6%	6	100.0%
223.0	61.2%	44	100.0%
2173.0	52.80%	912	100.0%
4035.6	86.80%	277	100.0%

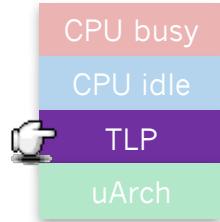
Legacy/improper OS designs
Performance overprovisioning

Anecdote
Voice UI

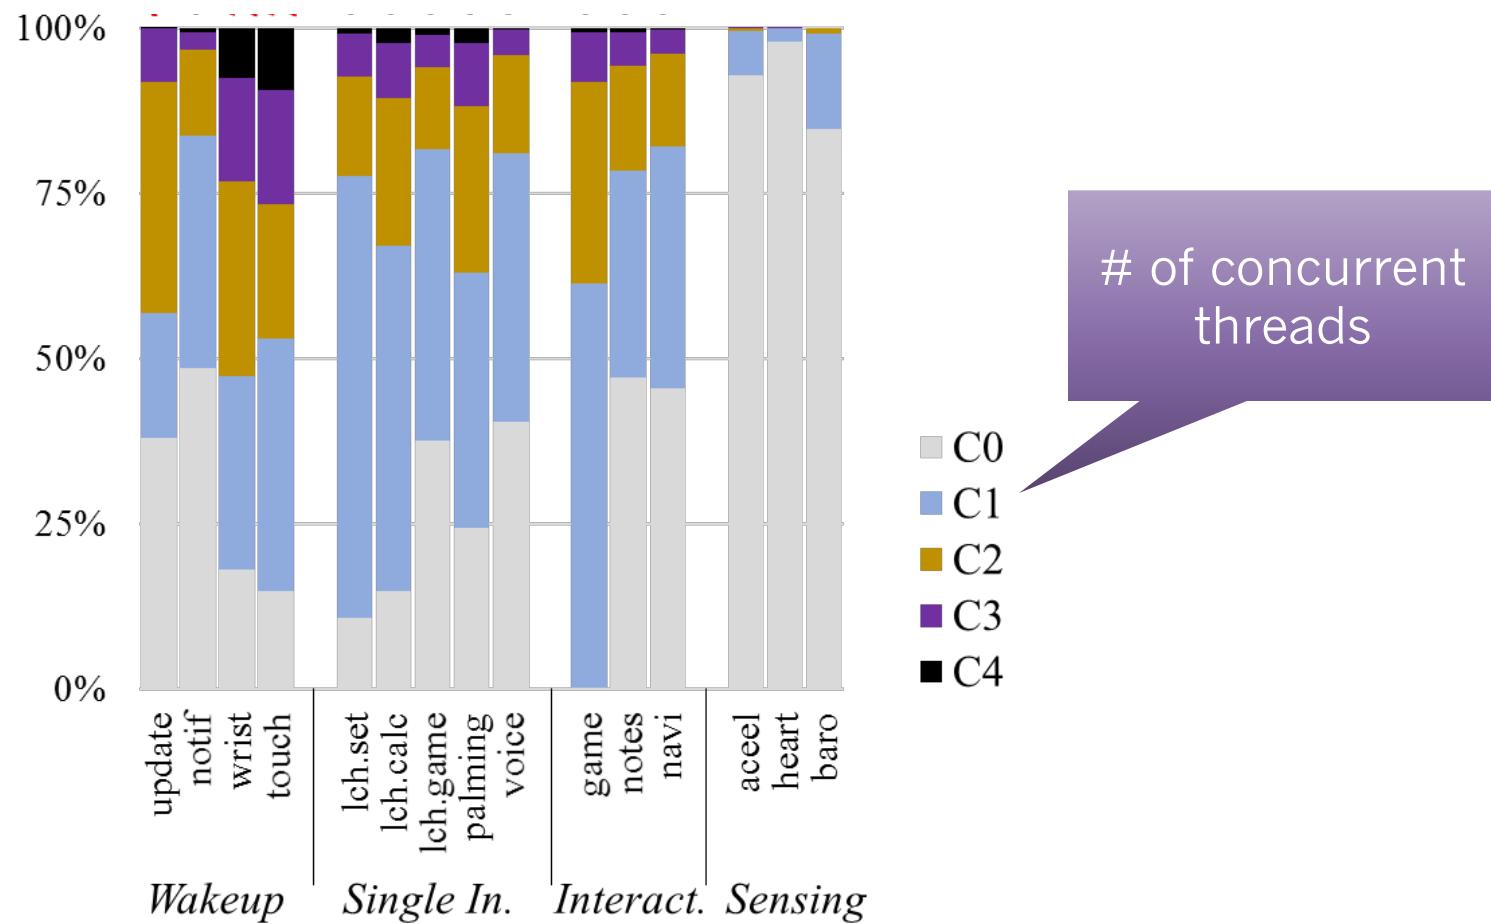


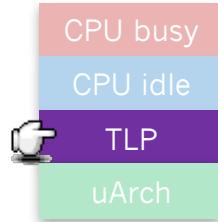
Substantial TLP on a par with desktop



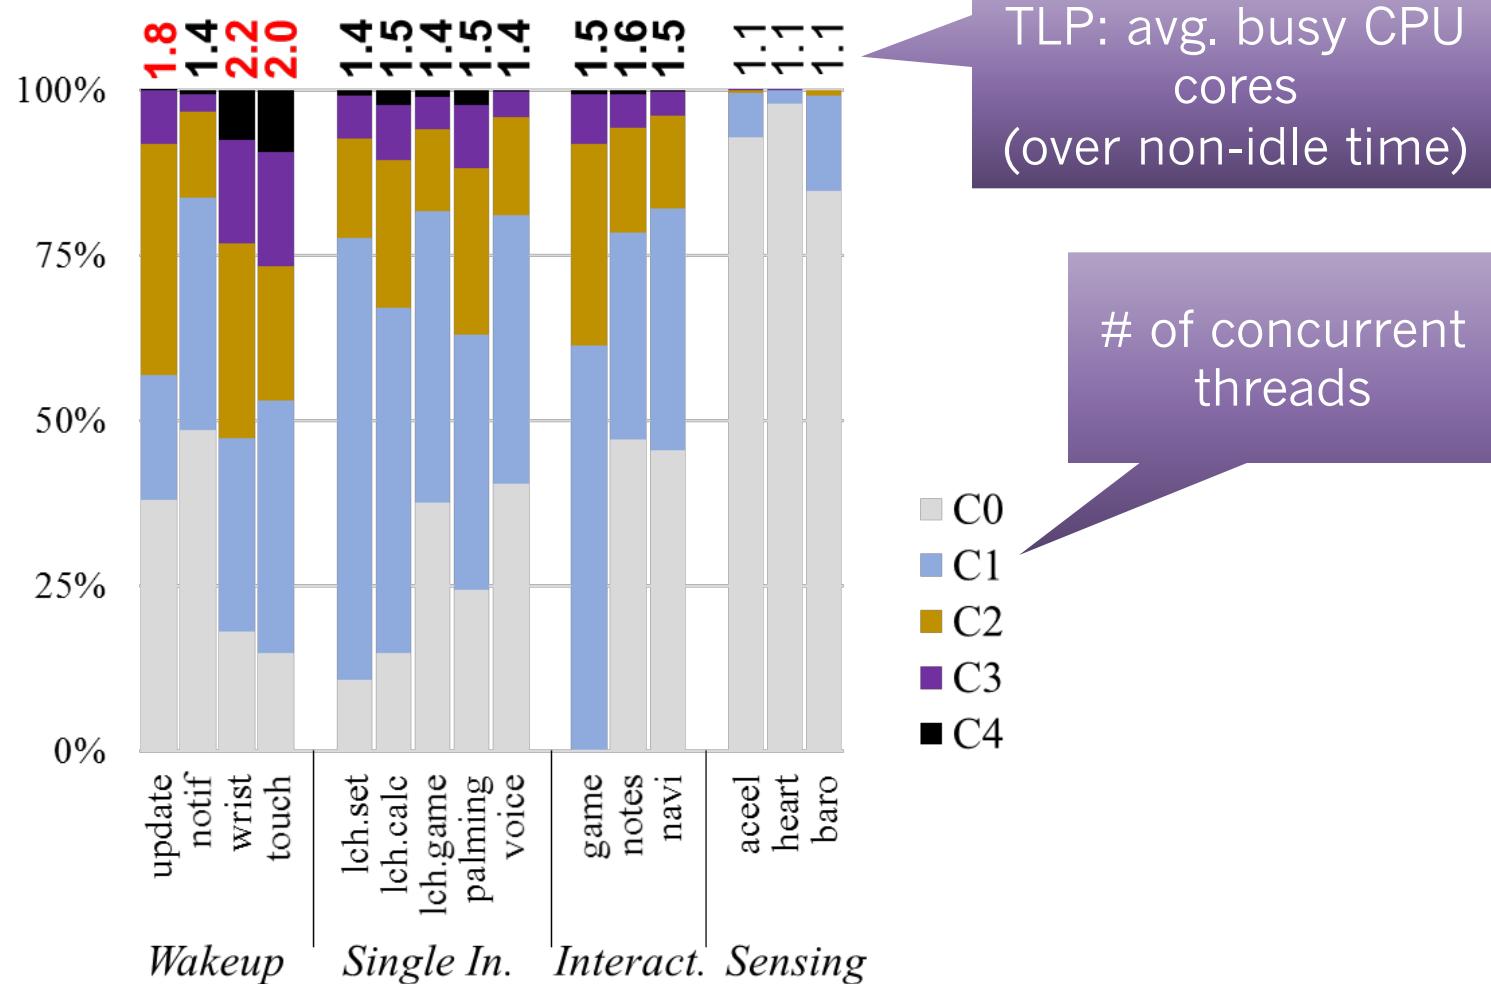


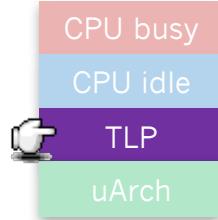
Substantial TLP on a par with desktop



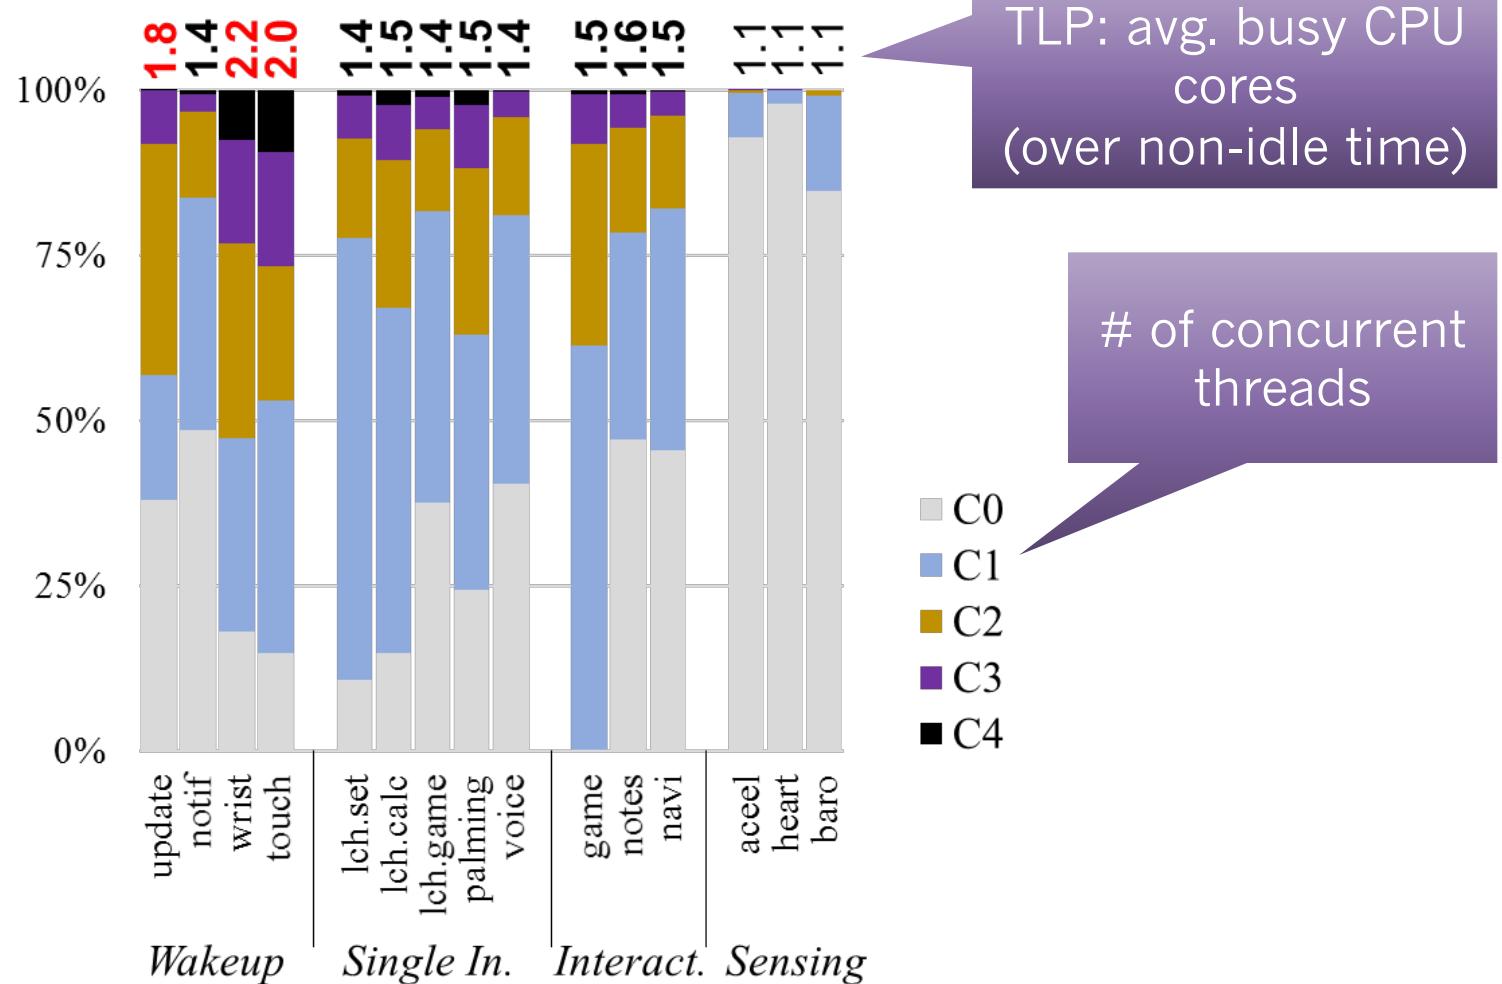


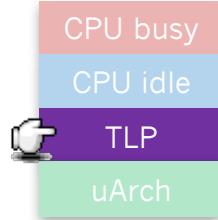
Substantial TLP on a par with desktop



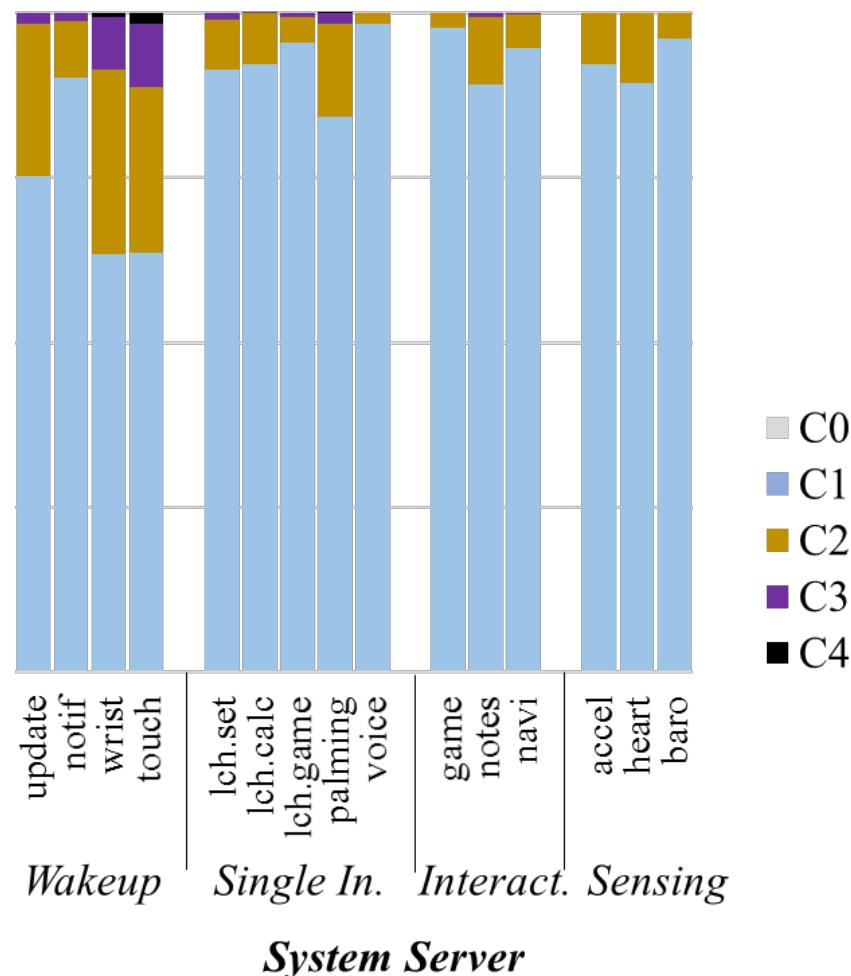


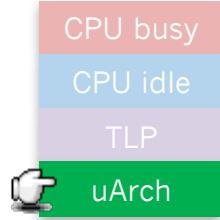
...due to short interactions.





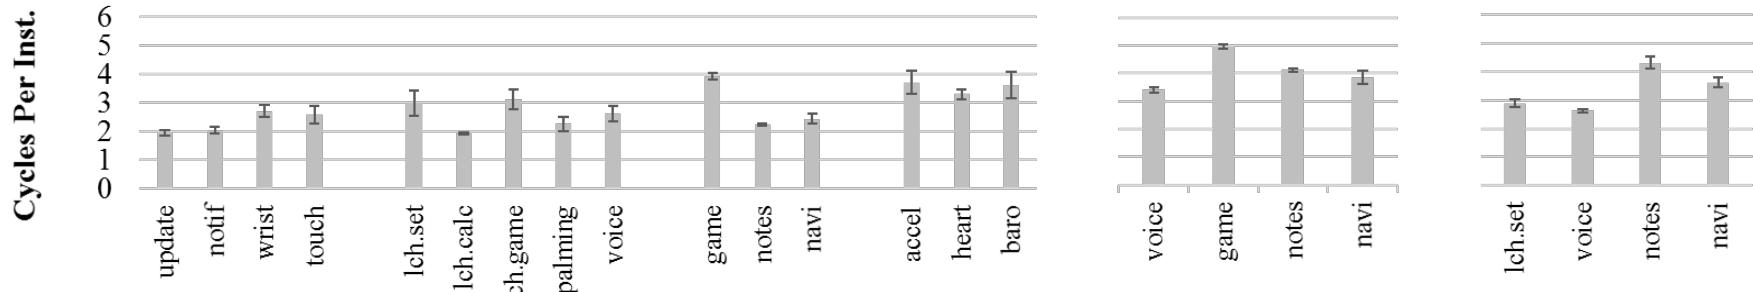
Apps are mostly single-threaded; OS contributes to TLP significantly.

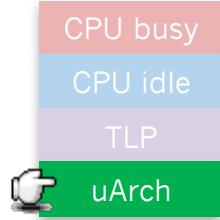




Wearable suffers from uArch inefficiency

Cycles-per-instruction (lower is better)
2 -- 5 (high!)





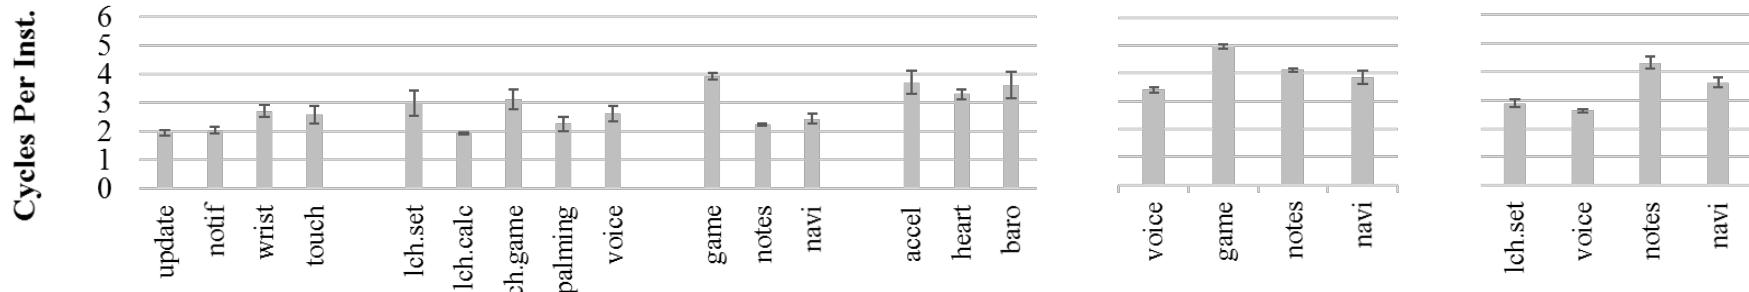
Wearable suffers from uArch inefficiency

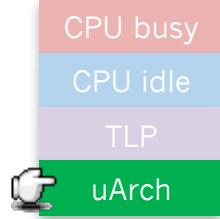
Cycles-per-instruction (lower is better)
2 -- 5 (high!)

Smartphone as a comparison

1.3 -- 2.5 web rendering

<2 SPEC INT





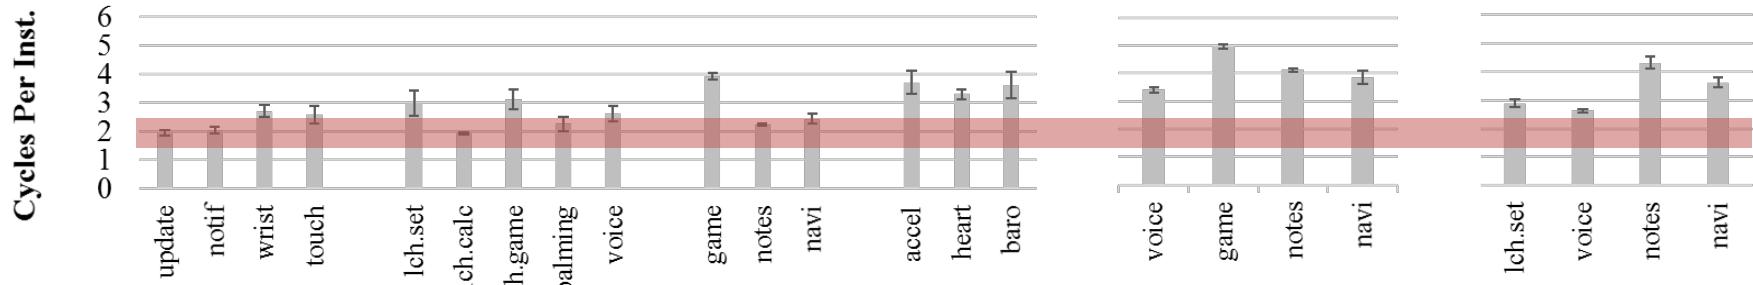
Wearable suffers from uArch inefficiency

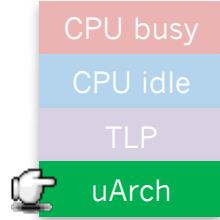
Cycles-per-instruction (lower is better)
2 -- 5 (high!)

Smartphone as a comparison

1.3 -- 2.5 web rendering

<2 SPEC INT





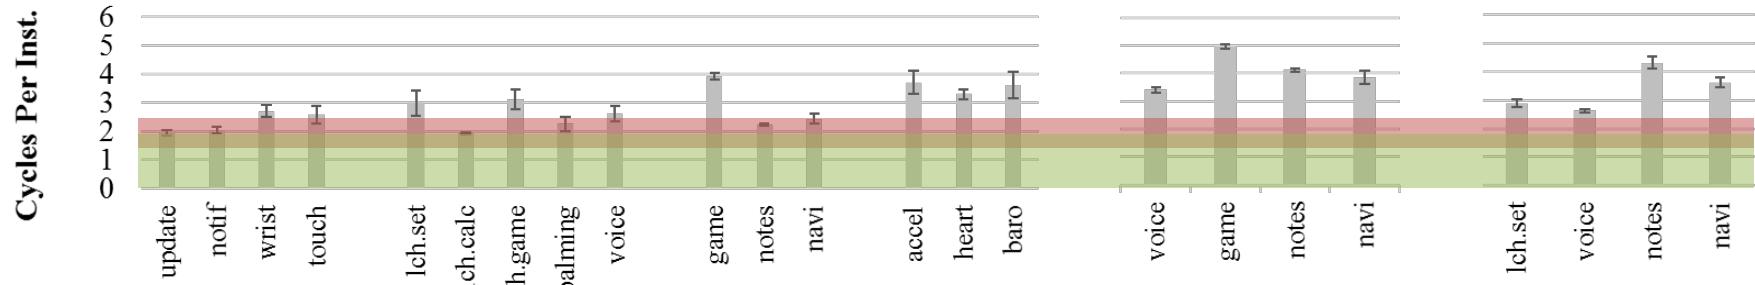
Wearable suffers from uArch inefficiency

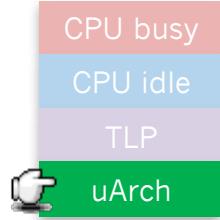
Cycles-per-instruction (lower is better)
2 -- 5 (high!)

Smartphone as a comparison

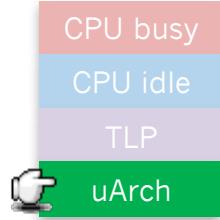
1.3 -- 2.5 web rendering

<2 SPEC INT





The major cause: complex OS code (L1 icache, iTLB, and branch predictor)



**The major cause: complex OS code
(L1 icache, iTLB, and branch predictor)**

**uArch problem will NOT be gone with
future wearable CPUs**

Four Aspects

CPU busy

- ◆ OS dominates ◆ Lots of cruft
- ◆ Skewed hot functions ◆ Legacy bottlenecks

CPU idle

- ◆ Anomalous ◆ OS flaws
- ◆ Too much performance

Thread-level parallelism

- ◆ Desktop-like ◆ OS-contributed

Microarchitectural behaviors

- ◆ Mismatch ◆ OS code complexity

Repair, don't overhaul (yet)

CPU busy

- ◆ OS dominates ◆ Lots of cruft
- ◆ Skewed hot functions ◆ Legacy bottlenecks

CPU idle

- ◆ Anomalous ◆ OS flaws
- ◆ Too much performance

Thread-level parallelism

- ◆ Desktop-like ◆ OS-contributed

Microarchitectural behaviors

- ◆ Mismatch ◆ OS code complexity

How about after that?
(i.e. “next-gen wearable OS”)

We probably will reach a point when OS
overhaul/redesign is justified.

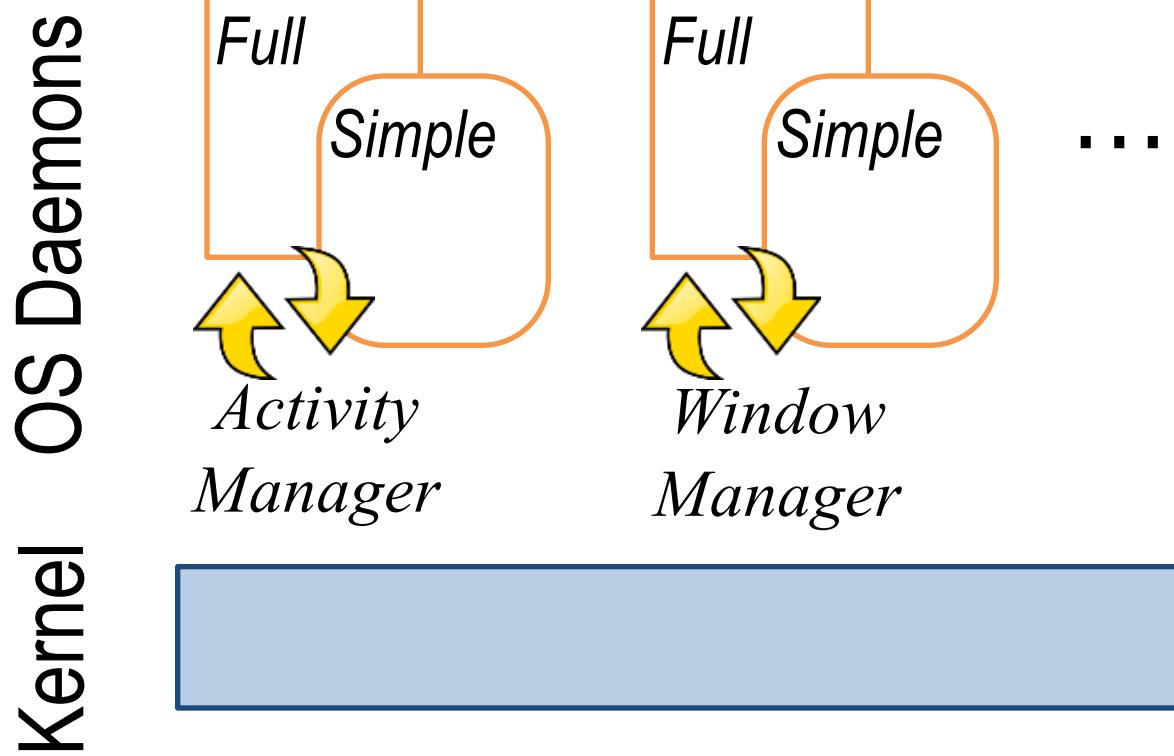


Specializing OS
for common, single-app scenarios

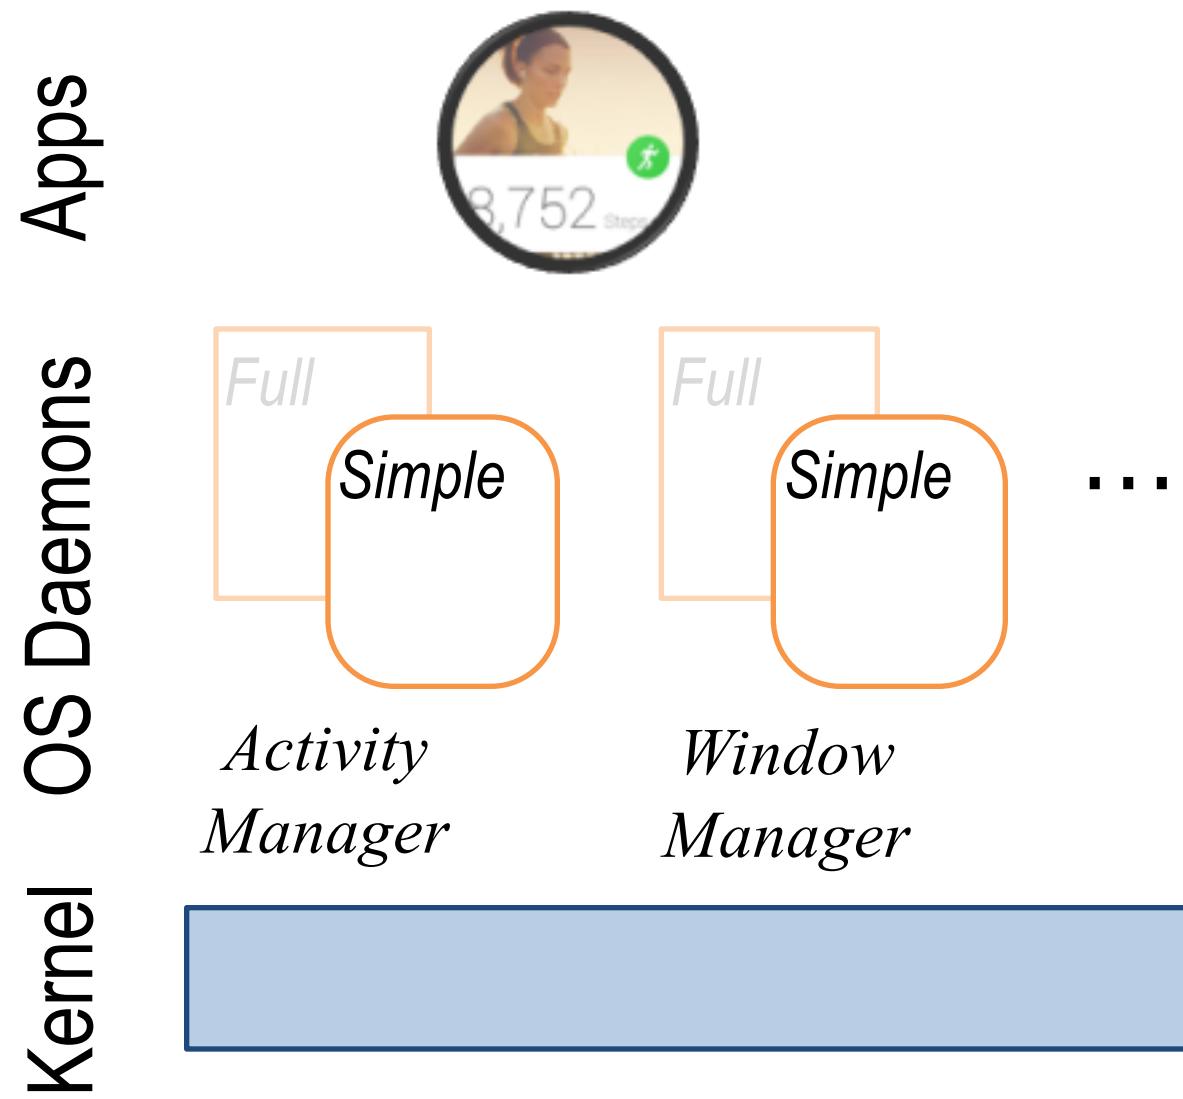
Restructuring OS for Wearable



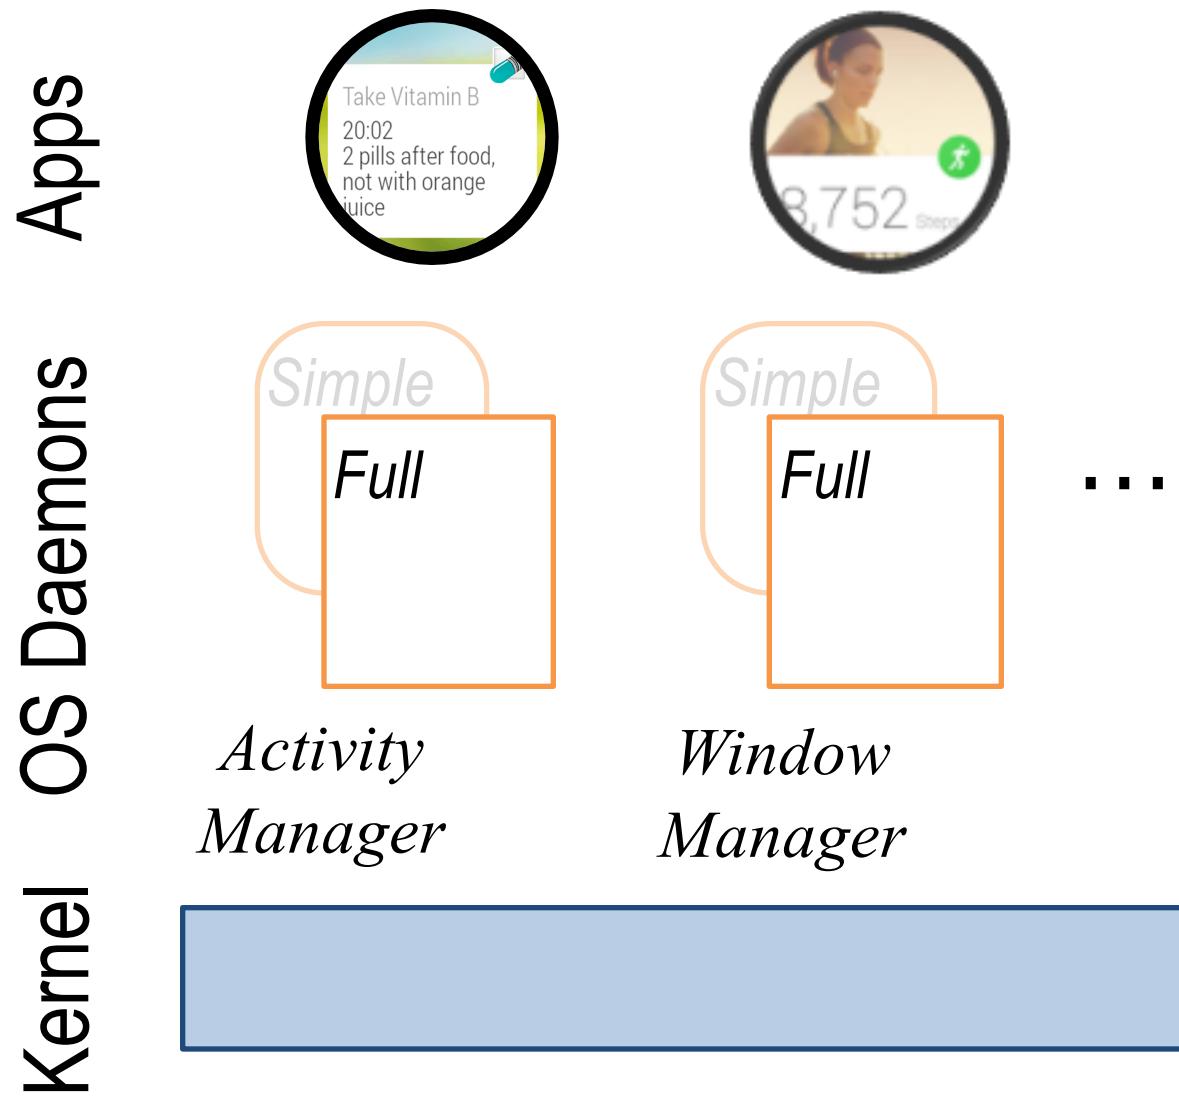
Specializing OS
for common, single-app scenarios



Restructuring OS for Wearable

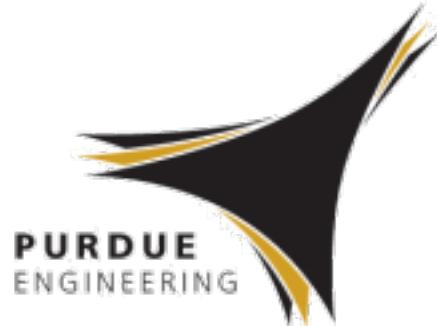


Restructuring OS for Wearable



Final takeaway

- Wearables: unique usage and hardware
- Many mobile OS tradeoffs are invalid
 - efficiency v.s. flexibility & programming ease
- Immediate actions: fixing individual OS components
- Future: OS specialization may be needed



Tools, data, and benchmark videos

xsel.rocks/p/wear

FAQ

- You forgot Apple Watch or Samsung Tizen.
- Isn't your discovery just some oversight of Google engineers?
- Aren't these things easy to fix?
- Doesn't multicore wearable sound crazy?
- Power! I want to learn about power.
- I bet the Android Wear team already fixed these!

Has Android Wear improved?

Android Wear 2.0 Developer Preview

New User Interface



- Material design for wearables
- Expanded notifications
- Darker UI

Standalone Apps



- Direct network access to cloud
- Apps run on watch even when your phone (Android or iOS) isn't with you

Watch Face



Complications API:
any watch face can
show data from
any app

Messaging



New input methods:
handwriting, keyboard,
Smart Reply, and 3rd
party IMEs

Fitness



Google Fit platform:
automatic activity
recognition and
data API

